# Shiv Chandekar

*Backend Engineer – Distributed Systems & Fintech Payments*

Mumbai, MH shivchandekar16@gmail.com +91-7767975374 linkedin.com/in/shiv-chandekar
sscba.github.io

## Summary

Java backend engineer with 2.5 years building distributed, event-driven fintech systems at scale. Delivered 80% latency reduction, 18x throughput gain, and 70% cost savings on systems processing 70M+ transactions/day using Spring Boot, Kafka, Redis, and YugabyteDB.

## Education

**Indian Institute of Technology Bombay**, Mumbai *Aug 2018 – Apr 2023*
Bachelor of Technology in Metallurgical Engineering and Materials Science
*Relevant Coursework:* Data Structures, Algorithms, Database Systems, Software Engineering

## Professional Experience

**Software Engineer** - *MindGate Solutions Pvt. Ltd., Mumbai* *Jul 2023 – Jan 2026*

**UPI Reconciliation System – Large-Scale Payment Processing Platform**

- Architected a Java/Spring Boot distributed reconciliation system processing **70M+ UPI transactions/day** across 10 settlement cycles, with near-real-time discrepancy detection across Switch, NPCI, and CBS
- Built Kafka ingestion pipelines (**20+ partitions, 5K–20K+ events/sec**) into YugabyteDB (Cassandra API) with a write-optimized schema for high-throughput
- Engineered a Redis Streams + Redis Gears reconciliation engine running **2M+ checks/cycle**, cutting cycle latency by **80%** and enabling real-time financial exception flagging
- Designed distributed job orchestration and Quartz pipelines across **6+ microservices** using Redis, reducing manual reconciliation effort by **60%**

**HyperStream Analytics POC – Cost-Optimized Transaction Processing**

- Migrated **70M+ UPI records** from Spark-Kafka to DuckDB vectorized processing, cutting batch runtime by **97% (4+ hours to 15 minutes)** with zero data integrity loss
- Replaced a Spark-kafka microservices architecture with a single-node DuckDB setup, achieving **18x throughput** and **70% cost reduction**

## Key Projects

**IntelliQuery Dashboard – AI-Powered Analytics Platform** *May 2025*

- Built a Llama3 + LangChain + DuckDB platform for natural language-to-SQL querying, reducing analyst turnaround time by **90% (60+ min to under 10 min)**
- Developed a LangChain RCA agent over error logs and system metrics, cutting mean time to diagnosis (MTTD) to **under 10 minutes** for production incidents
- Implemented JWT-based RBAC across **3+ data sources** with real-time Streamlit visualization, enabling self-serve analytics for non-technical stakeholders

**TaskManagerAPI – Enterprise-Grade Task Management System** *Jan 2026*

- Built a Spring Boot REST API (**20+ endpoints**) with JWT auth, RBAC, rate limiting (100 req/min), and email verification, sustaining **under 200ms** response time under load
- Cut critical-path latency by **90% (5–6s to 500ms)** via Spring `@Async` email processing, BCrypt tuning, and targeted database indexing
- Achieved **80%+ unit test coverage** with JUnit/Mockito and **34 Postman integration tests**; documented via OpenAPI/Swagger with Actuator observability

## Technical Skills

**Programming Languages:** Java, Python, C++, Lua, Bash
**Backend Technologies:** Spring Boot, Spring MVC, RESTful APIs, Microservices Architecture, JWT, RBAC
**Databases & Storage:** Oracle SQL, Redis, Yugabyte PostgreSQL, Yugabyte Cassandra, DuckDB, MongoDB
**Big Data & Streaming:** Apache Kafka, Apache Spark, Data Pipelines, Redis Streams
**Cloud & DevOps:** Docker, Kubernetes, Linux, Git, Maven
**Frontend & AI Technologies:** Streamlit, Llama3 LLM, LangChain
**Testing & Documentation:** Postman, Swagger, OpenAPI, JUnit, Mockito
**Concepts:** Distributed Systems, Event-Driven Architecture, System Design, Concurrency, High Availability, Agile